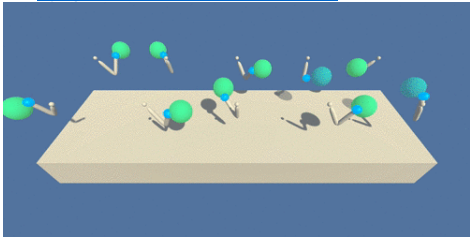
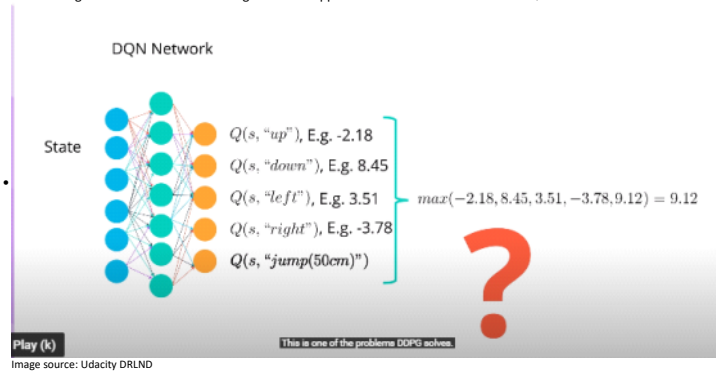
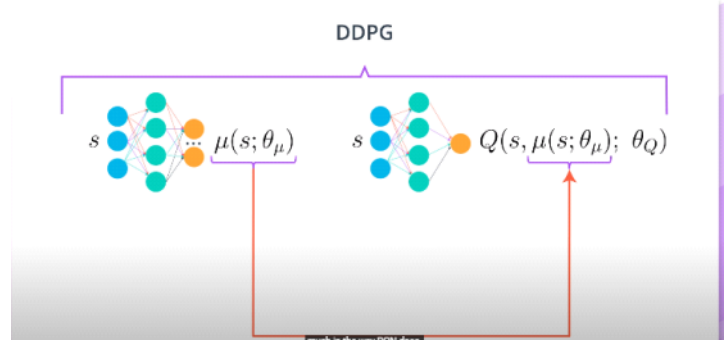


Project 2 - Continuous Control -DDPG

Sunday, 19 July 2020 11:39 PM

S.No	Topic	Details
1.	Project Goal	<p>This Reacher project is as part of Udacity Nanodegree - AI Deep Reinforcement Learning Expert and aims to develop an AI Agent - "a double-jointed arm" - move to target location in Continuous space using Policy-based 'Actor-critic' Methods using Deep Neural Networks.</p> <p>From <https://github.com/SENC/AIReacher/blob/master/README.md></p> 
2.	Scope	<ul style="list-style-type: none">• Develop an AI Agent using 'actor-critic' methods - which should learn the best policy to maximize its rewards by taking best actions in the given continuous environment• Goal The environment is considered solved, when the average (over 100 episodes) of those average scores is at least +30.• Decided to solve the First VersionOption 1: The task is episodic and the Agent must get an average score of +30 over 100 consecutive episodes
3.	Purpose	<ul style="list-style-type: none">• One of the primary goal of AI is to solve complex tasks in high dimensional , sensory inputs . Though Deep Q Network (DQN) proved to be high performance on many Atari video games but handles well in discrete and low-dimensional action spaces .DQN can't applied directly to continuous domain since the core part to find the action that maximizes the action-value function.• This project aims to build a model-free, off-policy actor-critic [Deterministic Policy - action-value] algorithm using deep function approximators that can learn policies in continuous space• DDPG Paper: https://arxiv.org/abs/1509.02971
4.	Solution Approach -Policy based Methods	<ul style="list-style-type: none">• Policy Gradients - An alternative to the familiar DQN (Value based method) and aims to make it perform well in continuous state space. Off-policy algorithm - Essential to learn in mini-batches rather than Online• Develop 'Actor-Critic' agent uses Function approximation to learn a policy (action) and value function<ul style="list-style-type: none">• Have 2 Neural Networks<ul style="list-style-type: none">◦ One for an Actor - Takes states information as an input and actions distribution as an output<ul style="list-style-type: none">▪ Take the action to move to next state and check the reward (Experience)and using TD estimate of the reward to predict the Critic's estimate for the next state◦ Next one for a Critic - Takes states as input and state value function of Policy as output.<ul style="list-style-type: none">▪ Learn to evaluate the state value function V_π using TD estimateTo calculate the advantage function and train the actor using this value.So ideally train the actor using the calculated advantages as a baseline.• Instead of having baseline using TD estimate , can use Bootstrapping to reduce the variance<ul style="list-style-type: none">◦ Bootstrapping - generalization of a TD and Monte-Carlo estimates<ul style="list-style-type: none">▪ TD is one step bootstrapping and MC is infinite bootstrapping▪ Mainly to reduce biasness and variances under controlled & fast convergence• Like DQN , have 'Replay Memory' - a digital memory to store past experiences and correlates set of actions -REINFORCE- to choose actions which mostly yields positive rewards<ul style="list-style-type: none">• Randomly collect experiences from the Replay Memory in to Mini-batches so the experiences may not be in same correlation as Replay Memory to train the Network successfully• Buffer size can be large so allowing the algorithm to benefit from learning across a set of uncorrelated transitions• Little change in 'Actor-critic' when using DDPG - to approximate the maximizer over the Q value of next state instead of baseline to train the value  <p>#1 In Actor NN : used to approximate the maximizer - an optimal best policy (action) deterministically - so the Critic learns to evaluate the optimal policy - Action-Value function for the best action</p> <p>Approximate the Maximizer - to calculate the new target value for training the action value function</p> $Q(s, \mu(s; \theta_\mu); \theta_Q)$  <p>Image source: Udacity DRLND</p> <p>Regular/local network - UpToDate network since training is done here but target network used to predict the stabilize strain</p> <p>#2 Soft Target updates:</p> <p>Weight of the target network are updated by having them slowly track the learned networks to improve the stability of learning</p>
5.	Algorithm	<p>Deep deterministic Policy Gradient</p> <p><u>Published as a conference paper at ICLR 2016</u></p>

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .
 Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$.
 Initialize replay buffer R
for episode = 1, M **do**
 Initialize a random process \mathcal{N} for action exploration
 Receive initial observation state s_1
for $t = 1, T$ **do**
 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise
 Execute action a_t and observe reward r_t and observe new state s_{t+1}
 Store transition (s_t, a_t, r_t, s_{t+1}) in R
 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R
 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|\theta^{Q'}$
 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$
 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{a=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for
end for

#Crux of DDPG in 9 simple steps for both AI and Human Values

6. Hyper parameters

```

BUFFER_SIZE = 500000 # replay buffer size
BATCH_SIZE = 256      # minibatch size
GAMMA = 0.997         # discount factor
TAU = 0.0013          # for soft update of target parameters
LR_ACTOR = 0.00013    # learning rate of the actor
LR_CRITIC = 1e-3       # learning rate of the critic
WEIGHT_DECAY = 0      # L2 weight decay

theta=0.17            # Noise Sampling
sigma=0.24
Random seed = 3       #reacherAI = AiAgent(state_size,action_size,random_seed=3)

```

7. Rewards

After multiple iteration and parameter tuning , finally Environment getting solved and please find the Score and reward chart

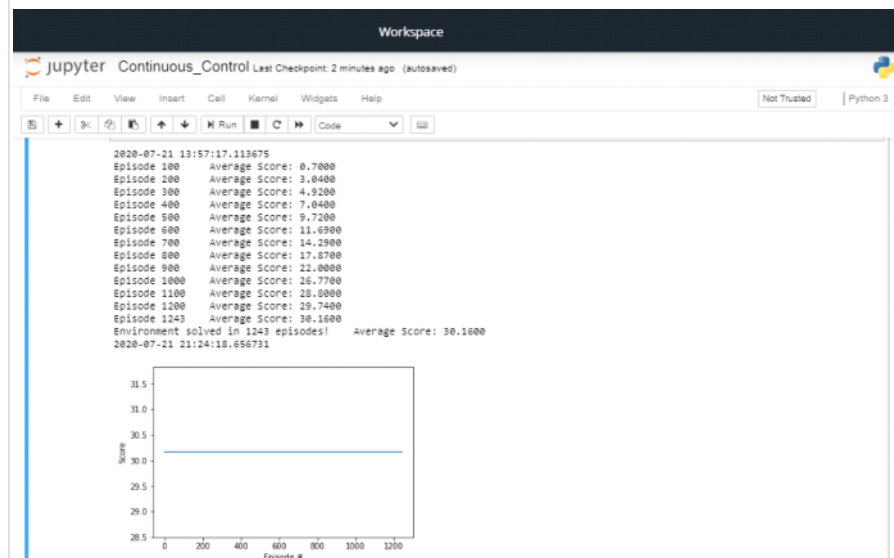
```

2020-07-21 13:57:17.113675
Episode 100   Average Score: 0.7000
Episode 200   Average Score: 3.0400
Episode 300   Average Score: 4.9200
Episode 400   Average Score: 7.0400
Episode 500   Average Score: 9.7200
Episode 600   Average Score: 11.6900
Episode 700   Average Score: 14.2900
Episode 800   Average Score: 17.8700
Episode 900   Average Score: 22.0000
Episode 1000  Average Score: 26.7700
Episode 1100  Average Score: 28.8000
Episode 1200  Average Score: 29.7400
Episode 1243  Average Score: 30.1600

```

Environment solved in 1243 episodes! Average Score: 30.1600

2020-07-21 21:24:18.656731



8. Source code

1. Nn_model.py -Convolutional Neural Network model with 3 layer architecture

		<div>#Test to create the instance of AiAgent</div> <pre>reacherAI = AiAgent(state_size,action_size,random_seed=9) print(reacherAI.actor_local) print(reacherAI.critic_local)</pre> <div>Actor((fc1): Linear(in_features=33, out_features=24, bias=True) (fc2): Linear(in_features=24, out_features=48, bias=True) (fc3): Linear(in_features=48, out_features=4, bias=True)) Critic((fcs1): Linear(in_features=33, out_features=24, bias=True) (fc2): Linear(in_features=28, out_features=48, bias=True) (fc3): Linear(in_features=48, out_features=1, bias=True))</div> <div>2. Agent : Agent with properties for</div> <ul style="list-style-type: none">• local and Target networks ,• soft update ,• Noise for exploration• Replay Memory for Experience Replay• step, act, reset , learn functions <div>3. Continuous Control.ipynb - Python Notebook covers all the Code and Results</div> <div>4. Checkpoint_actor30.pth - saved model weights for Actor</div> <div>5. Checkpoint_critic30.pth - saved model weights for Critic</div>
9.	Ideas for future work	<div>1. Work on Option 2 and try parallel learning PPO</div> <div>2. Solve a more difficult continuous control environment where the goal is to teach a creature with four legs to walk forward without falling.</div> <div>Ref https://github.com/Unity-Technologies/ml-agents/blob/master/docs/Learning-Environment-Examples.md#crawler</div>
10.	In Simple	THANKS TO UDACITY TEAM!!